

Sleep Quality Prediction Using Digital Biomarkers and Artificial Neural Networks

Course	BTech CSE — Machine Learning Project (2024–25)
Paper	Lee H, Cho M, Lee SW, Park S (2025). Predicting sleep quality with digital biomarkers and artificial neural networks. Frontiers in Psychiatry 16:1591448.
Members	Shreyas Borkar (202351132) Shabbir Anees (202352331) Vamshi Charan (202352337) Siddhikesh Gavit (202351040)
Tool	Google Colab (T4 GPU) · Python · TensorFlow · SHAP · LIME

This document is a complete reference guide for our ML project. It explains every concept, term, method, result, and pipeline in simple language — designed so that all team members can understand, revise, and present confidently. A presentation script for all 4 members is included at the end.

Contents

1. Project Overview & Aim
2. Problem Statement — All Parts
3. Research Paper Summary
4. Dataset — Structure, Features, Statistics
5. Data Preprocessing Pipeline
6. Feature Concepts Explained Simply
7. Machine Learning Pipeline
8. Neural Network Architectures — Detailed
9. Training Process & Loss Function
10. Leakage — What It Is and How We Fixed It
11. Results — Qualitative Analysis
12. Results — Quantitative Analysis
13. Explainability: LIME & SHAP
14. Bugs Found and Fixed (All 10)
15. Our Contributions Beyond the Paper
16. Presentation Script — All 4 Members

1. Project Overview & Aim

Modern life is stressful. Millions of people sleep poorly without knowing why — until it's too late. This project asks: **Can we look at what your body does *during the day* and predict how well you'll sleep *tonight*?**

We replicate and extend a 2025 research paper (Lee et al.) that collected heart-rate data from smartwatches worn by 82 people over 4 weeks. The paper found that a feature called the **LF/HF ratio** — a measure of your body's stress response — can predict whether you'll wake up during the night.

What We Predict

We predict **WASO** — **Wake After Sleep Onset**: the total minutes you spend awake after initially falling asleep. This is one of the most reliable markers of poor sleep quality.

We turn this into a **binary classification** problem:

- **Class 0** = No awakening (good sleep)
- **Class 1** = Had awakening episodes (disrupted sleep)

2. Problem Statement — All Parts

GIVEN

33–50% of adults experience insomnia symptoms (AASM data). Existing sleep assessment is either expensive lab-based polysomnography or subjective self-report diaries. Wearables like smartwatches collect continuous HRV, steps, gyroscope, and light data throughout the day.

OBJECTIVE

To predict whether a person will experience sleep disruption ($WASO > 0$) on the following night, using 7 days of wearable sensor data and questionnaire features. We compare 7 ML models and explain predictions using LIME and SHAP.

CONSTRAINTS

The original paper's raw dataset is not publicly available. We use a synthetic dataset that exactly mirrors the paper's statistics (Table 2). WASO labels are self-reported rather than measured by actigraphy, introducing potential bias. The cohort is young (avg 25 years) and not diverse.

APPROACH / TECHNOLOGY

Heart Rate Variability (HRV) features (LF/HF ratio, RMSSD) extracted from PPG sensor data using HeartPy library. 7-day sliding window sequences fed into LSTM, GRU, TCN, and Transformer models (all neural networks). Compared with ARIMA, Random Forest, and XGBoost (non-neural baselines). Platform: Google Colab T4 GPU. Stack: Python, TensorFlow/Keras, scikit-learn, SHAP, LIME, statsmodels.

MOTIVATION

Proactive healthcare: predict a bad sleep night before it happens, so interventions (stress reduction, bedtime routines) can be made in advance. The LF/HF ratio measured during daytime wakefulness is a proxy for stress load and has a known biological connection to sleep quality.

APPLICATION

Smartwatch apps that alert users if today's stress levels predict tonight's poor sleep. Clinical decision support for insomnia treatment planning. Personalized sleep coaching systems.

3. Research Paper Summary

Citation: Lee H, Cho M, Lee SW, Park S (2025). "Predicting sleep quality with digital biomarkers and artificial neural networks." Frontiers in Psychiatry 16:1591448.

What the paper did

Study design	Two experiments: Winter 2023 (Jan–Feb) and Summer 2023 (Jun–Jul). 28 days and 26 days respectively.
Participants	82 people (45 male, 37 female). Avg age 25.3 years. Mostly university students.
Device	Samsung Galaxy Watch Active 2 worn during daytime only (removed at night for charging). PPG sensor at 10 Hz, HR at 1 Hz, Steps at 1 Hz.
HRV extraction	HeartPy Python library extracted RMSSD and LF/HF ratio from 5-minute windows of PPG.
Questionnaires	PHQ-9 (depression), GAD-7 (anxiety), ISI (insomnia severity), KNHANES (stress), WHOQOL-BREF (quality of life). Collected weekly or biweekly.
Target	WASO (Wake After Sleep Onset) — self-reported daily in a sleep diary. Mean = 13.1 ± 6.5 min.
Best result	LSTM: Accuracy 0.904, Precision 0.913, Recall 0.899, AUROC 0.901.
Key finding	LF/HF ratio is the strongest predictor. ISI and WHOQOL-BREF also significant.
Split used	Random 80/20 split — SAME PARTICIPANT can appear in both train and test. This inflates the results (leakage issue we fixed in our project).

■ *Why we improved on it:* The paper's random split allows the model to see the same person in both training and testing. A person's sleep patterns are consistent, so the model just memorises individual patterns rather than learning general rules. Our subject-wise split completely separates people between train and test.

4. Dataset — Structure, Features, Statistics

Why Synthetic?

The original paper's raw data has not been released publicly. We created a synthetic dataset that **exactly mirrors Table 2 statistics** from the paper (means, standard deviations, class balance). We used a **latent stress variable** to create realistic biological correlations — meaning features are related to each other in the same way they are in real human bodies.

Dataset Statistics

Participants	82 (41 winter + 41 summer)
Days per person	27 on average
Total records	2,214 daily rows
Sequences (windows)	1,722 (7-day sliding window, stride=1)
Class 0 (no awak.)	41.5% of sequences
Class 1 (awakening)	58.5% of sequences
WASO stats	Mean = 13.1 min, Std = 6.5 min, Range = 0–38 min
Train subjects	66 participants (80%)
Test subjects	16 participants (20%) — NEVER seen during training

Features Collected

lf_hf (LF/HF Ratio)	Low Frequency to High Frequency power ratio of HRV. THE main predictor. High value = sympathetic nervous system dominant = stressed. Paper mean = 0.6 ± 0.2
rmssd	Root Mean Square of Successive Differences. Measures parasympathetic (calming) nervous system. High = relaxed and recovered. Paper mean = 135.4 ± 37.9 ms
steps	Daily cumulative step count. More steps = slightly better sleep. Paper mean = 4,954
accel	Magnitude of accelerometer reading — intensity of body movements
gyro	Gyroscope magnitude — rotational movement patterns
isi	Insomnia Severity Index questionnaire score (0–28). Higher = worse insomnia. Collected biweekly. Paper mean = 5.8 ± 4.2
whoqol	WHO Quality of Life score (WHOQOL-BREF). Higher = better life quality. Collected at baseline. Paper mean = 99.4 ± 13.3

Features removed: HR (heart rate — too correlated with rmssd) and Distance (too correlated with steps). Removal justified by Pearson correlation analysis and VIF > 5.

5. Data Preprocessing Pipeline

Step	What it does	Why it matters
1. KNN Imputation (k=3)	Fill ~3% missing sensor values using 3 nearest neighbors	Wearable sensors occasionally fail or have gaps. KNN preserves temporal structure.
2. Standard Scaler	Subtract mean, divide by std for each feature	Neural networks train faster and more stably when all features are on similar scales.
3. Pearson + VIF Filter	Remove features with $ r > 0.2$ with each other or VIF > 5	Multicollinearity (features too similar) confuses models. Keeping only independent features.
4. 7-Day Sliding Window	Create sequences: days t-6 to t → predict day t+1	Captures circadian rhythm & long-term trends. Matches paper's sliding window approach.
5. Class Weighting	Set class_weight = {0: 1.0, 1: neg/pos_ratio}	Class imbalance (negatives more common (58.5%)). Without weighting, model would ignore minority class.
6. ROC-Optimal Threshold	Find threshold T that maximises Youden's J	Default threshold (0.5) gave 70% sensitivity. Optimal threshold improved performance.

6. Feature Concepts Explained Simply

Heart Rate Variability (HRV) — The Core Signal

Your heart doesn't beat like a perfect metronome. The tiny variations in time between each heartbeat are called **Heart Rate Variability (HRV)**. These variations reveal the state of your **Autonomic Nervous System (ANS)** — the system that controls involuntary body functions.

- **High HRV** = body is relaxed, adaptable, in parasympathetic (rest-and-digest) mode
- **Low HRV** = body is stressed, in sympathetic (fight-or-flight) mode

How HRV is Measured from a Smartwatch

The Samsung Galaxy Watch uses a **PPG (Photoplethysmography) sensor** on the wrist. It shines green light on the skin and measures how blood absorbs it. Each heartbeat changes the blood volume, creating a wave pattern. The HeartPy library then extracts the R-R intervals (time between peaks) and computes RMSSD and LF/HF.

RMSSD — What It Measures

RMSSD = Root Mean Square of Successive Differences

Formula: $\sqrt{\text{mean}((RR[i+1] - RR[i])^2)}$

Think of it as: 'how much does the gap between heartbeats change from one beat to the next?'

- **High RMSSD** = good parasympathetic activity = calm, recovered
- **Low RMSSD** = stressed, fatigued

LF/HF Ratio — The Stress Meter

HRV signal is broken into frequency bands using Fourier transform:

- **LF (Low Frequency: 0.04–0.15 Hz)** = sympathetic nervous system activity (stress)
 - **HF (High Frequency: 0.15–0.40 Hz)** = parasympathetic nervous system (calm)
 - **LF/HF Ratio** = how much stress vs calm. Higher ratio = more stressed.
- Paper found: Higher LF/HF group had 14.9 min WASO vs 7.5 min for lower group ($p=0.012$).

ISI — Insomnia Severity Index

A 7-question self-report questionnaire (0–28 scale). It asks about difficulty falling asleep, staying asleep, early waking, how satisfied you are with your sleep, and how much it affects your daily functioning. Score 0–7 = no insomnia, 8–14 = subthreshold, 15–21 = moderate, 22–28 = severe.

WHOQOL-BREF — Quality of Life

World Health Organization Quality of Life Brief Version. 26 questions across 4 domains: physical health, psychological, social relationships, and environment. Higher score = better quality of life. It appears in our feature set because stress, QoL, and sleep are deeply interconnected.

WASO — The Target Variable

Wake After Sleep Onset (WASO) = total time (in minutes) spent awake after you first fall asleep until you finally wake up for good. It's a key marker of sleep *fragmentation*. Even if you spend 8 hours in

bed, 60 minutes of WASO means you only got 7 hours of actual sleep. The clinical cutoff for concern is WASO > 30 minutes.

7. Machine Learning Pipeline

Our complete pipeline goes from raw daily data to final predictions and explanations:

→ **Raw Data**

82 participants × 27 days = 2,214 daily records. Each record has 7 scaled features + binary WASO label.

→ **KNN Imputation**

Fill missing sensor values using 3 nearest neighbor rows.

→ **Standard Scaler**

Z-score normalise all 7 features. FIT ONLY on train.

→ **7-Day Window Builder**

Slide a window of 7 consecutive days per person. Each window = input X (shape: 7 × 7). Label = WASO on day 8.

→ **Subject-Wise Split**

66 people → train (1,386 sequences). 16 people → test (336 sequences). ZERO overlap.

→ **Class Weighting**

Weight = neg_count / pos_count given to model during training.

→ **Train 7 Models**

ARIMA, Random Forest, XGBoost, GRU, TCN, Transformer, LSTM.

→ **ROC Threshold**

For each model, find Youden's J optimal threshold on train ROC curve.

→ **Evaluation**

Accuracy, Precision, Recall, AUROC, Binary Cross-Entropy Loss.

→ **Explainability**

LIME (local per-prediction), SHAP (global across all predictions), Permutation Importance (feature shuffling test).

The 7-day window captures the weekly cycle in human sleep behavior. The paper specifically chose 7 days because people have distinct weekday vs weekend patterns. Using stride=1 means: days 1–7 predict day 8, days 2–8 predict day 9, etc.

8. Neural Network Architectures — Detailed

Why Neural Networks for Time Series?

Sleep quality depends on patterns across multiple days — not just today's stress level. Traditional models (ARIMA, Random Forest) treat each day independently or flatten all 7 days into one long feature vector, losing the temporal order. Neural networks, especially LSTM and GRU, are designed to **remember patterns across time**.

LSTM — Long Short-Term Memory (Best Model)

Architecture: BiLSTM(128) → LSTM(64) → LSTM(32) → BN → Dense(128→64→32→1)

LSTM is a type of Recurrent Neural Network (RNN) that solves the 'vanishing gradient' problem. It uses three gates to control what information flows through the network:

- **Forget Gate:** $f = \sigma(W_f \cdot [h, x] + b_f)$ — decides what to discard from memory
- **Input Gate:** $i = \sigma(W_i \cdot [h, x] + b_i)$ — decides what new information to add
- **Output Gate:** $o = \sigma(W_o \cdot [h, x] + b_o)$ — decides what to output as the hidden state
- **Cell Update:** $C = f \cdot C_{prev} + i \cdot \tanh(W_C \cdot [h, x] + b_C)$

Our upgrade — Bidirectional LSTM: We add a Bidirectional wrapper to the first layer. This means the sequence is read both forwards (day 1→7) AND backwards (day 7→1). For sleep prediction, this means day 3's HRV can be informed by what happens on day 6. This doubled the first layer's output from 128 to 256 units.

GRU — Gated Recurrent Unit

Architecture: GRU(128) → GRU(64) → GRU(32) → BN → Dense(64,32,1)

GRU is a simplified version of LSTM. It uses only 2 gates instead of 3:

- **Update Gate:** decides how much of the past to carry forward
- **Reset Gate:** decides how much of the past to forget

GRU is faster to train and uses less memory. In our results, GRU achieves the highest Precision (0.951) — meaning when it predicts a bad night, it's almost always right.

TCN — Temporal Convolutional Network

Architecture: Dense(64) → ResBlock(dilation=1) → ResBlock(dilation=2) → ResBlock(dilation=4) → GAP → Dense(32,1)

TCN uses **dilated causal convolutions** instead of recurrence. 'Causal' means the network can only look at past values, not future ones. 'Dilated' means gaps are inserted between filter elements, allowing the network to see a wider time range without more parameters. Dilation rates 1, 2, 4 give a receptive field of 7 days with just 3 layers. Residual (skip) connections help gradients flow during training.

Transformer — Self-Attention Network

Architecture: Dense(64) → PositionalEncoding → 3×EncoderBlock(4 heads, d=64) → GAP → Dense(64,32,1)

The Transformer uses **multi-head self-attention**: each day in the 7-day sequence can directly attend to (be influenced by) every other day, regardless of how far apart they are. We added **Positional Encoding** (not in the paper) so the model knows the day ordering. 4 attention heads let the model learn 4 different types of day-to-day relationships simultaneously.

9. Training Process & Loss Function

Loss Function — Binary Cross-Entropy

Since WASO is binary (0 or 1), we use **Binary Cross-Entropy loss**:

$$L = -[y \cdot \log(p) + (1-y) \cdot \log(1-p)]$$

Where y = true label (0 or 1) and p = model's predicted probability. When $y=1$ and the model predicts $p=0.9$, loss = $-\log(0.9) = 0.105$ (low, good). When $y=1$ and model predicts $p=0.1$, loss = $-\log(0.1) = 2.3$ (high, bad). The optimizer minimizes this loss by adjusting model weights.

Optimizer — Adam

Adam (Adaptive Moment Estimation) adjusts the learning rate for each parameter individually. It combines Momentum (smooths gradient updates) and RMSprop (scales by recent gradient magnitude). We use learning rate = $5e-4$ for LSTM. Adam handles the noisy gradients common in time-series data.

EarlyStopping

Training stops when **val_auc** (AUROC on the 15% validation split) doesn't improve for 15 epochs. The best weights are restored. This prevents overfitting — the classic symptom of a model that memorises training data but fails on new data.

AUROC — Area Under the ROC Curve

The ROC (Receiver Operating Characteristic) curve plots True Positive Rate vs False Positive Rate at every possible threshold. AUROC is the area under this curve.

- AUROC = 0.5 means random guessing
- AUROC = 1.0 means perfect prediction
- AUROC = 0.856 (our LSTM) means: if you pick a random 'awakening' person and a random 'no awakening' person, the model ranks the awakening person higher 85.6% of the time.

Key Metrics Explained

Accuracy	Correct predictions / Total predictions. Easy to understand but misleading on imbalanced data.
Precision	Of all predicted 'awakening', how many were actually awakening? ($TP / (TP+FP)$)
Recall	Of all actual awakenings, how many did we catch? ($TP / (TP+FN)$). Medical: 'sensitivity'.
AUROC	Ranking quality score. Best overall metric for imbalanced binary classification.
Loss	Binary Cross-Entropy. Lower = better. How far off are the probability predictions.

10. Data Leakage — What It Is and How We Fixed It

What Is Data Leakage?

Data leakage happens when information that the model should NOT have access to during training accidentally gets included. The model learns to cheat rather than learning the real pattern. The result: training looks great, but the model fails completely on truly new data.

Version 1 Problem (AUROC 0.92 — too high)	Random 80/20 split. The SAME PERSON appears in both train and test. Model memorises that 'person 42 always has awakening' rather than learning from features. This is called Subject Leakage.
Version 2 Problem (AUROC 0.55 — near random)	Fixed the split but now all inter-feature correlations were near-zero ($r \approx 0.02$). The synthetic data had no biological structure — features were sampled independently. The model had nothing coherent to learn.
Version 3 Fix (AUROC 0.87 — honest)	Subject-wise split (zero overlap) + latent stress variable creating realistic correlations. ROC-optimal threshold. Class weighting. All correct.

The 4 Leakage Tests We Ran (ChatGPT Recommended)

- **Shuffle-Label Test:** Randomly shuffle the WASO labels and retrain. If the model still gets high accuracy → leakage. Our result: shuffled AUROC ≈ 0.50 ■ (chance level, confirming no leakage)
- **Feature Ablation:** Remove ISI and WHOQOL (questionnaire features) and retrain sensor-only. If performance barely changes → questionnaires added minimal leak. Our result: AUROC dropped ~ 0.05 — healthy contribution
- **Subject-Wise Split:** Use completely different people in test. Our result: Strict separation with 66 train / 16 test
- **Permutation Importance:** Shuffle one feature at a time. If accuracy drops a lot → that feature is genuinely important. Our result: whoqol and isi drop accuracy most when shuffled → real signal

11. Results — Qualitative Analysis

EDA Plots (Exploratory Data Analysis)

LF/HF Group vs WASO (Figure 4 Replica)

We split participants into Lower LF/HF ($\leq$ median = 0.58) and Higher LF/HF groups. Wilcoxon Rank-Sum test showed $p=0.000$, confirming statistically significant difference. Higher LF/HF $\rightarrow$ more WASO (more disrupted sleep). This directly replicates Figure 4 of the paper.

Scatter Plot: LF/HF vs WASO ($r=0.744$)

Strong positive correlation in our data. Paper reported $r=0.22$ — our synthetic data has stronger signal because we explicitly built it in via the latent stress variable.

WASO by Day of Week

Weekend WASO is higher than weekday WASO. This matches the paper's supplementary Figure S5 and justifies why we include 7 full days (covering both weekday and weekend patterns).

LIME Analysis (Local Interpretable Model-Agnostic Explanations)

LIME explains why the model made a specific prediction for one particular sample. It works by perturbing (randomly changing) the input features and seeing which changes affect the prediction most. We ran LIME on 50 test samples (25 per class).

- Top features (our results): $rmssd > lf_hf > isi > accel > whoqol$
 - Paper reported: $lf_hf > ISI > WHOQOL$
 - Interpretation: Both HRV metrics ($rmssd$ and lf_hf) are in the top 2 — consistent because they both measure autonomic nervous system balance from the same PPG signal.

SHAP Analysis (Our Contribution)

SHAP (SHapley Additive exPlanations) is based on game theory. It assigns each feature a 'Shapley value' — how much it contributed to pushing the prediction away from the average. Unlike LIME (local only), SHAP gives both **global** (all samples) and **local** (individual sample) explanations consistently.

- SHAP top features: $isi > whoqol > lf_hf > rmssd$
 - Beeswarm plot shows: high isi (pink/red dots) $\rightarrow$ positive SHAP $\rightarrow$ pushes prediction toward awakening
 - The paper only used LIME. Adding SHAP is our original contribution.

12. Results — Quantitative Analysis

Complete Results Table

Model	Type	Accuracy	Precision	Recall	AUROC	Loss
Ensemble (LSTM+XGB) ★Hybrid		0.809	0.858	0.855	0.870	0.469
XGBoost	Tree	0.782	0.881	0.780	0.866	0.450
LSTM (best NN)	Neural Net	0.718	0.923	0.630	0.856	0.574
Transformer	Neural Net	0.747	0.877	0.723	0.854	0.474
TCN	Neural Net	0.771	0.840	0.811	0.854	0.494
GRU	Neural Net	0.709	0.951	0.604	0.850	0.607
Random Forest	Tree	0.762	0.865	0.762	0.845	0.443
ARIMA	Statistical	0.613	0.705	0.756	0.485	1.374

Key Observations

- Best AUROC: Ensemble(LSTM+XGB) = 0.870 — our novel addition not in the paper
- Best Precision: GRU = 0.951 — when GRU says 'bad night', it's right 95% of the time
- Best Recall: Ensemble = 0.855 — catches 85.5% of all actual awakening nights
- ARIMA (AUROC=0.485) is barely better than random guessing — pure statistics fail at temporal patterns
- All 4 neural networks achieve AUROC > 0.850 — confirming deep learning's superiority on this task
- Paper (random split) claimed LSTM AUROC = 0.901. Our honest split gives 0.856 — still excellent but lower due to no leakage

Confusion Matrix — LSTM Analysis

True Positive (143)	Correctly predicted awakening — model found the pattern → caught 63% of all real awakenings
True Negative (101)	Correctly predicted no awakening — 89% specificity for the no-awakening class
False Negative (84)	Missed actual awakenings — main weakness, lowering recall to 0.630
False Positive (12)	Said 'awakening' when there was none — very few, hence precision = 0.923
Threshold: 0.548	Youden's J (TPR-FPR maximised at this point). Default 0.5 gave all-class-1 predictions.

13. Explainability: LIME & SHAP

Black-box models (especially deep learning) are hard to trust in medical/health contexts. We need to explain WHY the model predicts a bad sleep night. We use two complementary approaches:

LIME (Local Interpretable Model-Agnostic Explanations)

What it does	Explains one specific prediction at a time. Creates small perturbations of the input and fits a simple linear model to understand local behavior.
How we used it	50 test samples (25 per class), 200 perturbations each. Aggregated importance across 7 features × 7 days.
Paper connection	Paper used LIME to produce Figure 5. Our LIME replicates this for our model.
Results	rmssd ranked #1, followed by lf_hf and isi. All HRV-related — biologically valid.
Limitation	LIME is local only — you need many runs to get a global picture. Can be unstable.

SHAP (SHapley Additive exPlanations) — Our Contribution

What it does	Assigns each feature a Shapley value — its fair contribution to the prediction, based on game theory. Provides both local (per-sample) AND global (average) explanations.
How we used it	TreeExplainer on Random Forest (exact, fast). Applied to test set last-day features.
Advantage over LIME	Globally consistent. Shapley values are provably fair. Doesn't need many samples.
Beeswarm plot	Each dot = one test sample. Red = high feature value. Positive SHAP = pushed toward class 1 (awakening). High isi + high lf_hf → strongly predicts awakening.
Results	isi ranked #1, followed by whoqol and lf_hf. Consistent with LIME's top 3.

Permutation Importance (Validation Test)

Shuffle each feature's values across test samples and measure how much accuracy drops. This confirms which features the model is genuinely relying on (not just correlated with). whoqol and isi showed the biggest drops — confirming they carry real predictive signal. lf_hf has high variance (wide error bar) — it's important but noisy.

14. Bugs Found and Fixed (All 10)

01 — CRITICAL: Subject Leakage

Problem Same person in train and test. Model memorised individuals → AUROC 0.92 (false).
m:

Fix: Subject-wise 80/20 split: 66 train / 16 test participants. Zero overlap enforced.

02 — CRITICAL: Zero Feature Signal

Problem Synthetic features sampled independently → $r \approx 0.02$ between features → model learned nothing
m: → AUROC 0.55.

Fix: Latent stress variable u per participant creates realistic biological correlations ($lf_hf \leftrightarrow rmssd$: $r = -0.79$, $lf_hf \leftrightarrow isi$: $r = +0.69$, etc.)

03 — HIGH: All-Class-1 Predictions

Problem Default threshold 0.5 with imbalanced data → model predicts everything as class 1. Confusion matrix: 15/110/10/205.
m:

Fix: ROC-optimal Youden's J threshold per model. LSTM optimal threshold: 0.548.

04 — HIGH: EarlyStopping Crash for GRU/TCN/Transformer

Problem GRU, TCN, Transformer compiled with `metrics=['accuracy']` only.
m: `EarlyStopping(monitor='val_auc')` crashes because 'auc' metric doesn't exist.

Fix: Added `tf.keras.metrics.AUC(name='auc')` to `compile()` for all 3 models.

05 — HIGH: SHAP Index Error

Problem `shap_vals[1]` assumes list. Newer SHAP returns 3D array → `IndexError` crash.
m:

Fix: Robust 3-branch check: `if isinstance(list) → [1], elif ndim==3 →[:, :, 1], else direct.`

06 — MEDIUM: LIME Feature Matching Broken

Problem String-in-string: 'isi' matches inside 'rmssd'. Wrong features getting importance.
m:

Fix: Pre-built dict `feat_name_to_base`. Exact `startswith()` lookup for feature attribution.

07 — MEDIUM: LSTM Overfitting

Problem Train AUROC 0.86 vs Val AUROC 0.72 — classic overfitting gap.
m:

Fix: Dropout 0.35→0.45, added L2 regularization ($1e-4$), monitor `val_auc` (not `val_loss`).

08 — MEDIUM: Permutation Importance Title Hardcoded

Problem Title said 'lf_hf top' regardless of actual top feature. Our results showed whoqol on top.
m:

Fix: Dynamic detection: `top_feature = pi_df.iloc[0]['feature']`. Title reads actual result.

09 — LOW: google.colab Crash Outside Colab

Problem: Cell 25 download cell crashes with `ModuleNotFoundError` when run in Jupyter/VS Code.

Fix: `try/except ImportError`: if Colab → `files.download()`, else → print file path.

10 — MEDIUM: Scaler Fit on Full Data

Problem: `StandardScaler.fit_transform()` on the entire dataset before splitting. Test set statistics contaminate training normalisation.

Fix: `Scaler.fit_transform(df_train)`, `Scaler.transform(df_test)`. Imputer same approach.

15. Our Contributions Beyond the Paper

1. Subject-Wise Split

Methodology [Methodology]

Paper used random 80/20 split allowing same participant in train and test. We enforce strict subject-wise separation: 66 train participants, 16 test participants, zero overlap. This is more rigorous, more honest, and a stronger academic methodology.

2. Bidirectional LSTM

Architecture [Architecture]

Paper used unidirectional LSTM. We wrapped the first layer in Bidirectional(), giving the model access to both forward (day1→7) and backward (day7→1) temporal context. Architecture: BiLSTM(128)→LSTM(64)→LSTM(32)+BN+Dense(128→64→32→1).

3. SHAP Analysis

Explainability [Explainability]

Paper only used LIME (local explanations). We added SHAP (SHapley Additive exPlanations) using TreeExplainer on Random Forest. SHAP gives globally consistent feature importance based on game theory — theoretically superior to LIME.

4. LIME vs SHAP Comparison

Explainability [Explainability]

We produce a direct side-by-side feature ranking comparison between LIME (local) and SHAP (global). Both methods agree on top features (HRV + insomnia measures), validating the biological interpretation.

5. Leakage Audit Suite

Validation [Validation]

Added 3 explicit leakage tests: (a) shuffle-label test to confirm signal is real, (b) sensor-only feature ablation, (c) permutation importance. None of these appear in the original paper.

6. Ensemble Model (LSTM+XGBoost)

Modeling [Modeling]

Soft-voting ensemble combining LSTM probability (55%) and XGBoost probability (45%). Achieves the highest overall AUROC in our study = 0.870. Not present in the original paper at all.

7. ROC-Optimal Threshold Tuning

Evaluation [Evaluation]

Paper used fixed threshold 0.5 for all models. We compute the optimal decision threshold per model using Youden J statistic. This eliminates the all-class-1 prediction bias from class imbalance.

16. Presentation Script — All 4 Members

The presentation is divided into 4 parts. Each person speaks for approximately 3–4 minutes. Total time: ~15 minutes + 5 minutes Q&A.; Slides 1–8 map to the professor's template exactly.

■ SPEAKER 1 — Shreyas Borkar (202351132) | Slides 1 & 2

Slides: Title + Problem Statement

Good morning everyone. My name is Shreyas Borkar, and on behalf of our group — Shabbir, Vamshi, and Siddhikesh — I'd like to present our Machine Learning project.

[Point to slide title] Our project is titled: '*Sleep Quality Prediction Using Digital Biomarkers and Neural Networks*.' We replicated and extended a 2025 research paper published in Frontiers in Psychiatry.

[Move to Slide 2 — Problem Statement]

Let me start with the problem we are solving. According to the American Academy of Sleep Medicine, between 33 and 50 percent of adults experience insomnia symptoms. 6 to 10 percent have clinically significant insomnia — and poor sleep leads to cardiovascular disease, cognitive decline, and higher healthcare costs.

Currently, sleep is measured either in expensive labs — using polysomnography — or through self-report diaries which are subjective and unreliable. Both methods are reactive: they measure sleep after the fact. Our approach is proactive.

Our objective: using 7 days of wearable sensor data collected during the day, we predict whether a person will experience sleep disruption — measured as WASO, that is Wake After Sleep Onset — on the following night. We frame this as a binary classification: Class 0 means no awakening, Class 1 means disrupted sleep.

The motivation is straightforward: if your smartwatch can tell you this morning that tonight's sleep might be bad because your stress levels have been elevated — you can take action: reduce workload, exercise, sleep earlier. This is the future of personalized health management.

■ SPEAKER 2 — Shabbir Anees (202352331) | Slides 3 & 4

Slides: Dataset + Neural Architecture

Thank you Shreyas. I'm Shabbir Anees, and I'll explain our dataset and the neural network architectures we used.

[Move to Slide 3 — Dataset]

The research paper conducted two experiments in 2023 — one in winter, one in summer — with 82 participants wearing Samsung Galaxy Watch Active 2 devices during waking hours. Since the original dataset is not publicly released, we created a synthetic dataset that exactly mirrors the paper's Table 2 statistics.

Our key innovation here was using a latent stress variable — a hidden per-participant score that simultaneously drives LF/HF up, RMSSD down, ISI up, and WHOQOL down. This creates the realistic biological correlations between features that real data has. In our previous versions, features were

sampled independently and the model learned nothing.

We use 7 features: LF/HF ratio and RMSSD — both HRV metrics from the PPG sensor; steps and accelerometer and gyroscope — physical activity; and ISI and WHOQOL-BREF — from questionnaires. A 7-day sliding window creates 1,722 sequences. Critically, we use a subject-wise 80/20 split — 66 participants for training, 16 completely different participants for testing. This prevents subject leakage.

[Move to Slide 4 — Neural Architecture]

Now the neural networks. We implemented four neural architectures. GRU uses gated recurrent units with update and reset gates to capture temporal dependencies. TCN — Temporal Convolutional Network — uses dilated causal convolutions with dilation rates 1, 2, 4, giving a 7-day receptive field. Transformer uses multi-head self-attention with 4 heads and positional encoding — we added positional encoding, which the paper's version lacked. And LSTM — our best model — uses Bidirectional LSTM which reads sequences both forwards and backwards.

The LSTM stack is: BiLSTM(128) → LSTM(64) → LSTM(32) → BatchNorm → Dense(128) → Dropout(0.45) → Dense(64) → Dense(1, sigmoid). We use Binary Cross-Entropy loss, Adam optimizer at $5e-4$, and EarlyStopping monitoring `val_auc` with patience 15.

■ SPEAKER 3 — Vamshi Charan (202352337) | Slides 5 & 6

Slides: Result 1 (Qualitative) + Result 2 (Qualitative/Quantitative)

Thank you Shabbir. I'm Vamshi Charan, and I'll present our qualitative results and the explainability analysis.

[Move to Slide 5 — Result 1 Qualitative]

Our first result replicates Figure 4 from the paper. We split participants into those with lower LF/HF ratio and higher LF/HF ratio based on the median cutoff of 0.58. The Wilcoxon Rank-Sum test — used because both variables are non-normal — shows a statistically significant difference at p equals zero-point-zero-zero-zero.

What does this mean? LF/HF ratio is a measure of sympathetic nervous system activity. When you are stressed, LF frequency increases relative to HF frequency. Our results confirm that people with higher daytime stress — higher LF/HF — experience significantly more nighttime awakenings.

We also ran LIME on our LSTM model, which is the paper's main explainability method. LIME explains individual predictions by perturbing the input and fitting a local linear model. Our top features were `rmssd`, `lf_hf`, and `isi` — all HRV and stress-related. This replicates the paper's Figure 5 finding.

Our contribution beyond the paper: SHAP analysis. SHAP uses Shapley values from game theory to give globally consistent feature importance. The beeswarm plot on the right shows that high ISI scores — indicating severe insomnia symptoms — consistently push predictions toward 'awakening'. High LF/HF does the same. Both LIME and SHAP agree on the top 3 features, validating the biological interpretation.

[Move to Slide 6 — Result 2 Qualitative/Quantitative]

The confusion matrix for our LSTM at the optimal threshold of 0.548 shows: 143 true positives, 101 true negatives, 84 false negatives, and 12 false positives. The low false positive count of 12 explains the very high precision of 0.923 — when our model says 'bad sleep tonight', it is correct 92% of the time.

We also show permutation importance, which was recommended by ChatGPT as a leakage validation test. By shuffling each feature and measuring accuracy drop, we confirm that whoqol and isi carry genuine predictive signal. Negative bars for rmssd indicate it contributes through correlation with other features rather than independently — not a sign of leakage.

■ SPEAKER 4 — Siddhikesh Gavit (202351040) | Slides 7 & 8

Slides: Result 3 (Quantitative) + Conclusion & References

Thank you Vamshi. I'm Siddhikesh Gavit, and I'll present the full quantitative results and our conclusions.

[Move to Slide 7 — Result 3 Quantitative]

Slide 7 shows the complete model comparison across all 8 models — including our novel Ensemble model that wasn't in the paper. Looking at AUROC, which is the most reliable metric for imbalanced binary classification:

Our Ensemble of LSTM and XGBoost achieves the highest AUROC of 0.870 through soft-voting — 55% weight on LSTM's probability and 45% on XGBoost. XGBoost alone achieves 0.866. All four neural networks — LSTM at 0.856, Transformer at 0.854, TCN at 0.854, and GRU at 0.850 — significantly outperform ARIMA which only reaches 0.485, barely better than random guessing.

An important comparison: the paper reported LSTM AUROC of 0.904 using a random split. Our subject-wise split gives 0.856. The difference — 0.048 — is the cost of being honest. Our result is genuinely generalisable to new patients. The paper's result partially reflects subject memorisation.

[Move to Slide 8 — Conclusion]

To summarize our key conclusions: First, neural networks — especially LSTM — are superior to traditional time-series and tree-based models for this task, because they capture the temporal dependencies across 7 days. Second, honest evaluation methodology matters — subject-wise split gives genuinely generalizable results. Third, the LF/HF ratio is the core biomarker: daytime sympathetic nervous system activity reliably predicts nighttime sleep fragmentation — confirmed by LIME, SHAP, and statistical testing.

Our novel ensemble model achieves the best overall performance at AUROC 0.870. Our explainability work — combining LIME with SHAP — goes beyond the paper to provide both local and global feature importance. We also conducted a rigorous leakage audit with 4 tests, addressing all 10 bugs identified across three development iterations.

For future work, we would like to validate on a real-world diverse dataset such as the MESA Sleep Study, extend to multi-class WASO severity prediction, and explore on-device inference for real-time smartwatch deployment.

We are happy to take any questions. Thank you.

Quick Reference — Likely Q&A; Questions

Q: What is WASO?

Wake After Sleep Onset — total minutes awake after first falling asleep. Binary: 0=none, 1=any.

Q: Why LF/HF ratio?

Daytime sympathetic nervous system stress directly predicts nighttime sleep fragmentation. $r=0.74$ correlation.

Q: Why subject-wise split?

Random split lets the model see the same person in training and testing — it memorises individuals. Subject-wise split forces generalisation to new people.

Q: Why bidirectional LSTM?

Reading the 7-day sequence both forward and backward gives richer temporal context. Our upgrade over the paper.

Q: What is AUROC?

Area Under ROC Curve. How well model ranks awakenings above non-awakenings. 0.5=random, 1.0=perfect.

Q: Why SHAP over LIME?

SHAP is globally consistent (Shapley values from game theory). LIME is local only — can vary between runs.

Q: Why ensemble?

LSTM excels at temporal patterns, XGBoost at feature precision. Combining both (55/45) gives AUROC 0.870 — best result.

Q: What's the dataset?

Synthetic, mirroring paper's Table 2 statistics. 82 participants $\times$ 27 days. Latent stress variable creates biological correlations.

Quick Glossary of All Terms Used

WASO	Wake After Sleep Onset — minutes awake after initially falling asleep
HRV	Heart Rate Variability — tiny beat-to-beat time variations that reveal ANS state
LF/HF Ratio	Low/High Frequency HRV ratio — stress indicator. Higher = more stressed
RMSSD	Root Mean Square Successive Differences — HRV metric for parasympathetic recovery
PPG	Photoplethysmography — wrist light sensor that measures blood flow / heartbeat
ISI	Insomnia Severity Index — 7-question self-report scale (0–28)
WHOQOL-BREF	WHO Quality of Life questionnaire — 26 items, higher = better QoL
LSTM	Long Short-Term Memory — RNN with forget/input/output gates for sequential data
GRU	Gated Recurrent Unit — simpler RNN with update/reset gates
TCN	Temporal Convolutional Network — dilated causal convolutions for time series
Transformer	Attention-based NN — self-attention lets each step attend to all others
AUROC	Area Under ROC Curve — ranking quality score (0.5=random, 1.0=perfect)
LIME	Local Interpretable Model-Agnostic Explanations — per-prediction explainability
SHAP	SHapley Additive exPlanations — globally consistent feature importance (game theory)
Subject-Wise	Train/test split where different people are in each set — prevents memorisation
Leakage	When test information accidentally enters training — inflates performance
Youden's J	Metric = TPR – FPR, maximised to find optimal classification threshold
Binary CE	Binary Cross-Entropy — loss function for two-class classification
Adam	Adaptive Moment Estimation — optimizer combining momentum and RMSprop
VIF	Variance Inflation Factor — multicollinearity detector. VIF > 5 = remove feature
KNN Imputer	k-Nearest Neighbors imputation — fills missing values with average of k=3 neighbors